

Data Partitioning

• Why Partition?

- At some point, single node based DB isn't going to cut it: concurrent read/write traffic will go up as you scale, so latency & throughput are affected
- A no-SQL like solution could be better, but historical colbase & close cohesion w/ traditional DB make this expensive
- Data Partitioning: enables us to use multiple nodes where each node manages some part of the whole data

• Sharding (Partitioning)

- To divide load among multiple nodes, we need to partition the data using sharding
- The partitions or "shards" should be even in size approximately
- If the partitions are unbalanced, most queries will go to a few nodes
- Heavily loaded shards will form a bottleneck: hotspots
- Two ways to shard: vertical / horizontal partitioning

• Vertical Partitioning

- used to increase speed of data retrieval from a table consisting of columns w/ very wide text or a binary large object (blob)
 - ↳ In this case column w/ large text is split into diff table

• Horizontal Sharding

- Splits tables row-wise

↳ each partition of the original table distributed over DB servers is called a shard

- Two types: key-range based sharding & hash based

- Key-Range based

↳ each partition assigned a continuous range of keys

↳ sometimes, a DB consists of multiple tables bound by foreign key relationships. In such a case, the horizontal partition is made using the same key on all tables in a relation

- tables (or subtables) that belong to the same partition key are distributed to one DB shard

• Basic Design Techniques

- partition key in the customer mapping table

↳ this table resides on each shard, stores partition keys used in the shard

↳ Applications create a mapping logic b/t the partition keys & DB shards by reading this table from all shards to make the mapping efficient

- Partition key column is replicated in all other tables as a data isolation point

↳ It has a tradeoff: ↑ storage ↓ speed to find desired shard.

- Primary keys are unique across all DB shards

↳ avoids key collision during data migration among shards & merging of data in online analytical processing environments

• Hash Based Sharding

- uses hash like function on an attribute

- use hash function on the key to get a hash value & then mod by # partitions

- When we're found an appropriate function for keys, we may give each partition a range of hashes rather than a range of keys. Any key whose hash occurs inside that range will be kept in that partition
- Consistent Hashing
 - assigns each server or item in a distributed hash table to a place on an abstract circle, called a ring, irrespective of the num of servers in the table
 - ↳ permits servers & objects to scale w/out compromising system's performance

Pros:

- Easy to scale horizontally
- Increases throughput & improves latency

Cons:

- Randomly assigning nodes in the ring may cause non-uniform distribution

- Rebalance the Partitions
 - Query load can be imbalanced across nodes/table:
 - ↳ data distribution of the data isn't equal
 - ↳ Too much load on a single shard
 - ↳ There's an \uparrow in query traffic, we need to add more nodes to keep up

Solutions:

- Avoid Hash mod n
- Fixed num of partitions
 - ↳ # partitions created is fixed @ time we set up DB
 - ↳ create more partitions than nodes & assign them to nodes

- Dynamic Partitioning

- ↳ when the size of the partition reaches the threshold, it's split evenly into 2 partitions
- ↳ one of the 2 splits is assigned to one node & the other one to another node. In this way, the load is \div equally
- ↳ Cons: Difficult to apply dynamic rebalancing while serving reads/writes

- Partitioning & Secondary Indexes

- ↳ How do we partition if we have access to records thru secondary indices?
- Partition secondary indexes by document
 - ↳ each partition has its secondary indices covering just the docs in that partition
 - ↳ This can be expensive