

Databases

- Problem Statement: can we build software w/out using dbs?

- WhatsApp

- ↳ need place to store info: list of people's names, & messages associated w/ them

- need this storage to be permanent

- ↳ could use simple file storage, but it has limits

- Limitations of File Storage

- We can't offer concurrent management to separate users accessing the storage files from diff locations

- We can't grant different access rights to diff users

- How will the system scale for 1000s of entries?

- How will we search content for diff users in a short time?

- Solution: Databases

- Organized collection of data that can be managed & accessed easily

- DBs are created to make it easier to store, retrieve, modify, & delete data in connection w/ diff data processing procedures

- SQL (relational dbs)

- ↳ organized & have predetermined schemas: phone books w/ contact numbers & addresses

- NoSQL (non-relational dbs)

- ↳ file directories that store anything from contact info to shopping preferences: unstructured, scattered, dynamic schema

- Advantages
 - Manage large data
 - Retrieve accurate data
 - Easy to update (using data manipulation language)
 - Security: only authorized users can access the DB
 - Data Integrity: diff constraints on data
 - Availability: DBs can be replicated on diff servers, & can be concurrently updated. Replicas ensure availability
 - Scalability: DBs are divided (using partitioning) to manage the load on a single node

• Relational DBs (RDBs)

- adhere to particular schemas before storing data
- data stored in RDBs has prior structure
- This model organizes data into one or more relations (also called tables) w/ a unique key for each tuple (instance)
- Instances
 - ↳ stored in rows
 - ↳ attributes of instances stored in columns
 - ↳ each tuple has a unique key — a tuple in one table can be linked to a tuple in other tables by storing the primary keys in other tables (known as foreign keys)
- SQL used for manipulating data on the DB
- RDBs are simple, robust, flexible, & can scale well while managing generic data well
- RDBs provide ACID:
 - ↳ Atomicity: A transaction is considered an atomic unit, therefore, either all the statements within a transaction will successfully execute, or none will. If a statement fails, it should be aborted & rolled back
 - ↳ consistency: DB should be in a consistent state, & should remain in a consistent state after every transaction

- ex:) If multiple users are viewing a record from a DB, it should return a similar result each time
 - ↳ Isolation: Multiple transactions happening concurrently shouldn't affect each other
 - Final DB state should be the same as the transactions were executed sequentially
 - ↳ Durability: System should guarantee that completed transactions will survive permanently in the DB even in sys failure events
- DB management systems: MySQL, Oracle DB, Microsoft SQL Server, IBM DB2, Postgres, SQLite

• Why relational DBs?

- RDBs are the go-to for structured data
- flexibility: In the context of SQL, data definition language (DDL) gives us flexibility to mod the DB
- Reduced Redundancy: info related to a specific entity appears in one table while relevant data to that specific entity appears in other tables linked thru foreign keys
 - ↳ this process: normalization, additional benefit of removing an inconsistent dependency.
- Concurrency: Important factor while designing an enterprise DB
 - ↳ data is read & written by many users @ the same time
 - ↳ we need to coordinate these actions to avoid inconsistencies in data — for ex:) double booking of hotel rooms
 - ↳ This concurrency concept is handled thru transactional access to the data
 - ↳ Transactions are atomic ops: so it works in

error handling to either roll back or commit a transaction on successful execution.

- Integration: process of aggregating data from multiple sources

↳ commonly used: shared DB where multiple apps store data

↳ This way, all apps can easily access each other's data while the concurrency control measures handle the access of multiple apps

- Backup & Disaster Recovery

↳ Export & Import ops make backup & restoration easy

↳ cloud based DBs perform continuous mirroring to avoid data loss

• Drawback: Impedance Mismatch

- What is it? Diff b/w the relational model & in-memory data structures

- The relational model organizes data into a tabular structure w/ relations & tuples

- SQL ops on this structured data yields relations aligned w/ relational algebra

- It has limitations:

↳ values in a table take simple values that can't be a structure or list

↳ In-memory, complex data structures can be stored

↳ to make complex structures compatible w/ the relations, we would need a translation of the data in light of relational algebra

↳ So: Impedance mismatch requires translation b/w 2 representatives

• Non-relational databases (NoSQL)

- used in apps that require semi/un-structured data
- Low latency & flexible data models
- Simple Designs:
 - ↳ don't deal w/ impedance mismatch
- Horizontal Scaling:
 - ↳ NoSQL is preferred due to its ability to run DBs on a large cluster
 - ↳ If num users ↑ by a lot, NoSQL makes it easy to scale since data is stored on a single document instead of multiple tables over nodes
 - ↳ NoSQL DBs spread data across many nodes & balance queries automatically, node failure results in transparent replacement of the failed node without application disruption.
- Availability: node replacement performed w/out app downtime
- Support for unstructured & semistructured data
- Cost: NoSQL DBs are free

• Types of NoSQL databases

- Document DB: mongo DB, Google Cloud Firestore
 - ↳ Tree Structure: can include maps, collections, scalars
 - ↳ designed to work w/ XML, JSON, BSON, etc.
- Graph DB: Neo4J, Orient DB, Infinite Graph
 - ↳ Graph structure: nodes represent entities, edges rep. relationships b/t entities
 - ↳ use case: social apps
- Columnar DB: Cassandra, HBase, HyperTable, Amazon Simple DB
 - ↳ efficient for large num of aggregation & data analytics queries

↳ reduces amount of data required to load from the disk

- Key-Value DBs: Amazon DynamoDB, Redis, Memcached DB

↳ stores data in key-value pairs

↳ useful in session based apps (web apps, store user info in a DB during a session)

↳ unique ID for a user session for easy access

• Drawbacks of noSQL DBs

- lack of standardization: no "relational algebra" like w/ RDBS.

- Consistency

↳ we can't have strong data integrity (like primary & referential integrity in relational DBs)

↳ Data might not be strongly consistent but slowly converging using a weak model like eventual consistency

• Choosing the right DB

relational	non-relational
<ul style="list-style-type: none">- Structured Data- ACID is required- Size of data is small & can be stored on a node	<ul style="list-style-type: none">- Unstructured data- need to serialize & deserialize data- size of data is large