


Load Balancers

• What is it?

- To serve millions of incoming requests, thousands of servers (or even hundreds of thousands) work to share the load
- They share the load via load balancers
- The goal of a LB:
 - ↳ fairly divide all client's requests among the pool of available servers
 - done to avoid crashing the servers
- You may not need an LB for a service that entertains only a few hundred → thousand requests
- For those that need them, load balancers provide systems:
 - ↳ **scalability**: by adding servers, the capacity of the app can be ↑ seamlessly. They make upscaling or downscaling transparent to the end user
 - ↳ **Availability**: even if some servers go down, the system will still be available. LBs also are meant to hide faults & failures of the servers.
 - ↳ **Performance**: They quickly forward requests to servers w/ a lesser load so the user gets a faster response. This improves both performance & resource utilization



- Placing LB's
 - Usually, LB's sit b/t clients & servers
 - ↳ Requests go thru to servers & back to clients via the LB layer
 - Not the only point that they're used
 - Place LB's between:
 - ↳ end users of the application & web servers / application gateway
 - ↳ web servers & application servers that run the business / app logic
 - ↳ b/t app servers & db servers


• LB services

- **Health checking**: LBs use the **heartbeat protocol** to monitor the health (& thus, reliability) of end servers
- **TLS termination**: They reduce the burden on end servers by handling TLS termination w/ the client
- **Predictive analytics**: Can also predict traffic patterns thru analytics performed over traffic passing thru or using stats of traffic over time
- **Service discovery**: client's requests are forwarded to appropriate hosting servers by inquiring about service registry
- **Security**: They can mitigate DOS attacks @ diff layers of the OSI model (layers 3, 4, 7)

• Global server load balancing (GSLB)

- involves the distribution of traffic load across multiple geographical locations
- ensures that globally arriving traffic load is intelligently forwarded to a **data center (dc)**

- Makes forwarding decisions based on:
 - ↳ user's geo location
 - ↳ num of hosting servers in diff locations
 - ↳ health of DCs.
- offers auto zonal failover



- GSLB can forward requests to 3 diff data centers
- each LB layer w/in a DC will maintain info about the health of LBs & server farm
- GSLB uses this info to drive traffic decisions & forward traffic load based on each regions config & monitoring in
- Local load balancing
 - LB-ing achieved w/in a data center
 - this type of LB-ing focuses on improving efficiency & better resource use of the hosting servers in the data center
 - They behave like a reverse proxy

• Advanced Details

- Algorithms

- ↳ round robin scheduling: requests forwarded to a server pool in repeated, sequential manner
- ↳ weighted RR: servers w/ higher capability have ↑ weight
- ↳ Least connections: Nodes w/ fewer connections get requests
- ↳ Least response time
- ↳ IP Hash: Server decided by IP Hash
- ↳ URL Hash: Some services w/in the app are provided by specific servers only. URL Hash function is used

- Static vs. Dynamic Algorithms

↳ Static:

- don't consider the changing state of the servers
- task assignment done using existing knowledge on servers config
- not complex

↳ dynamic:

- consider recent state of servers
- maintain state thru communication (adds overhead)
- require diff LBs to communicate
- can be modular → more complexity, but better decisions

- Stateful vs. Stateless

↳ Stateful

- involves maintaining a state of the sessions est. b/t clients & hosting servers
- stateful LB includes state in its algo to perform LB-ing

↳ Stateless

- maintains no state: faster, lightweight
- use consistent hashing to make forwarding decisions

- not as resilient as stateful LBs: consistent hashing alone isn't enough to route a request to the correct app server

↳ therefore a local state may also be required

- state managed across diff LBs: stateful

- state managed w/in a LB: stateless

• Types of LBs

- Layer 4

↳ LB-ing performed on the basis of transport protocols like TCP UDP

↳ maintain connection/session w/clients

↳ ensure TCP/UDP communication ends up being forwarded to the same backend server

↳ TLS termination: some L4 LBs support it

- Layer 7

↳ based on data of app layer protocols

↳ possible to make app aware forwarding decisions based on HTTP headers, URLs, cookies, & other app specific data

— for example, User ID

↳ TLS termination, rate limiting users, HTTP routing, header rewriting

• LB deployment

- DNS's are tier 0 LBs

- ECMPs (equal cost multipath) routers are tier 1 LBs

↳ divides traffic based on IP or RR or weighted RR

↳ Tier 1 LBs balance load across diff paths to higher tier LBs

↳ play a vital role in horizontal scalability of higher tier LBs

- ↳ Tier 2 LBs: include layer 4 LBs
- Make sure that for any connection, all packets are forwarded to the same tier 3 LBs
 - Glue b/t tier 1 & tier 3 LBs.

- ↳ Tier 3 LBs: In direct contact w/ backend servers
- perform health monitoring of servers @ HTTP level
 - This tier enables scalability by evenly distributing requests among healthy back-end servers
 - High availability by health-monitoring servers directly
 - Idea: leave computation & data serving to app servers & effectively utilize LB commodity machines for trivial tasks